

# Actuarial Reserving System

## Technical Architecture Documentation

**Version:** 1.0

**Last Updated:** 2026-02-03

**Author:** Auto-generated from codebase analysis

**Last Updated:** 2026-02-03

**Author:** Auto-generated from codebase analysis

# 1. Project Structure

## 1.1 Directory Layout

```
reserving-agent/
    app.py                                # Streamlit entry point (landing page)
    pages/
        1_Reported_Claims.py
        2_Summary.py
        3_AI_Team.py
    src/
        agents/
            orchestrator.py
            methodology.py
            selection.py
            reserving.py
            validation.py
            reporting.py
            qa.py
            schemas.py
            llm_utils.py
            settings.py
            main.py
            chain_ladder.py
            data_loader.py
            enhanced_workflow.py
            visualizer.py
            stochastic_reserving/
                mack_model.py
                bootstrap.py
                alternative_methods/
                    bornhuettet_ferguson.py
                    cape_cod.py
                model_selection/
                    factor_estimators.py
                    model_selector.py
                    validation_framework.py
                    error_metrics.py
                    statistical_tests.py
                    diagnostics/
                        diagnostic_tests.py
                        volatility_analysis.py
                        residual_analysis.py
                    tail_fitting/
                        tail_estimator.py
                    ml_models/
                        gradient_boosting_factors.py # GBM factor selection
                        anomaly_detector.py       # Triangle anomaly detection
                scenario_analysis/
                    economic_scenario_generator.py
                    stress_testing.py
                    tail_risk.py
        data/
            processed/                      # Primary data files
                reported_absolute_losses.csv
                earned_premium.csv
            sample_triangles/               # Benchmark datasets
            outputs/                       # Generated reports
            requirements.txt
```

## 1.2 File Responsibilities

File	Responsibility	Key Classes/Functions
`orchestrator.py`	Request routing, workflow orchestration	`Orchestrator.route_request()`, `stream_workflow()`
`methodology.py`	Analysis type determination	`MethodologyAgent.plan_analysis()`
`selection.py`	LLM-driven method selection	`SelectionAgent.analyze_and_select()`
`reserving.py`	Actuarial calculations	`ReservingExecutionAgent.execute()`
`validation.py`	Result validation	`ValidationAgent.validate()`
`reporting.py`	Report generation	`ReportingAgent.generate_report()`
`qa.py`	Question answering	`QASpecialistAgent.answer_query()`
`schemas.py`	Data models	`ReservingOutput`, `ValidationReport`, etc.
`llm_utils.py`	OpenAI integration	`LLMClient.get_completion()`

## 1.3 Dependencies

```
# Core
numpy>=1.21.0
pandas>=1.3.0
pydantic>=2.0.0

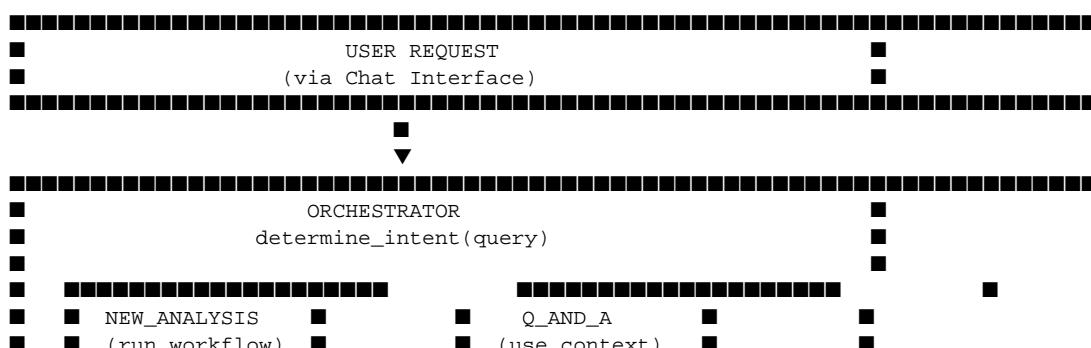
# UI
streamlit>=1.28.0
plotly>=5.0.0

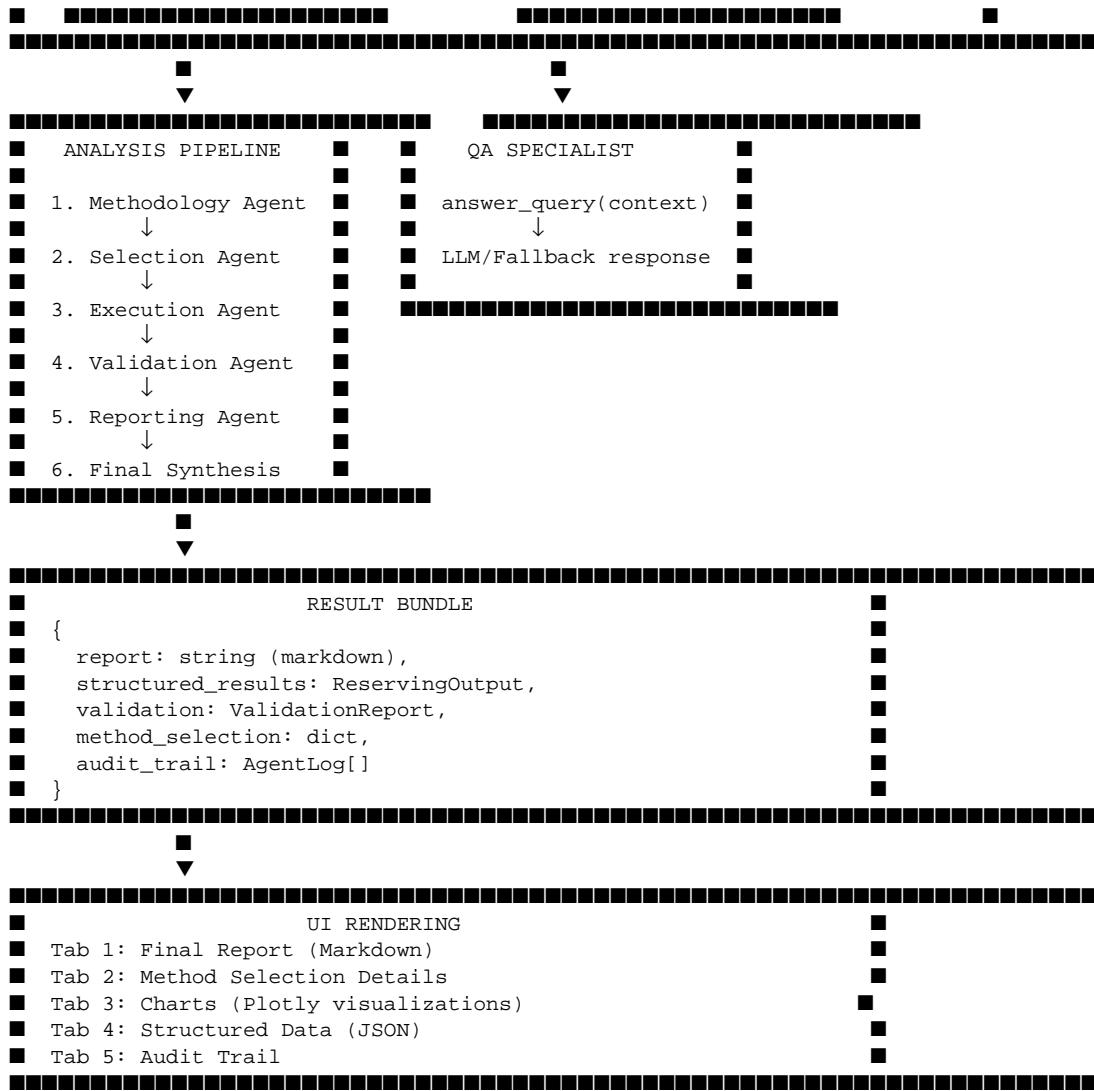
# Data
openpyxl>=3.0.0

# LLM
openai>=1.0.0
```

## 2. Core Workflow

### 2.1 Request Flow Diagram





## 2.2 Streaming Architecture

L'Orchestrator usa un **generatore Python** per streaming in tempo reale:

```
def route_request(self, message, current_result, inputs):
    intent = self.determine_intent(message, has_context)

    if intent == "NEW_ANALYSIS":
        yield {"step": "router", "message": "Starting analysis..."}

        for update in self.stream_workflow(message, inputs):
            yield update # Real-time status updates

    elif intent == "Q_AND_A":
        yield {"step": "router", "message": "Answering question..."}
        answer, log = self.qa_agent.answer_query(message, current_result)
        yield {"step": "qa", "status": "done", "data": answer}
```

## **Update Format:**

```
{  
    "step": "methodology" | "selection" | "execution" | "validation" | "reporting" | "complete",  
    "status": "running" | "done",  
    "message": str,  
    "data": Any  # Agent-specific output  
}
```

## 3. Agents Analysis

### 3.1 Orchestrator

\*\*File:\*\* `src/agents/orchestrator.py`

\*\*Classe:\*\* `Orchestrator`

#### Responsabilità:

- Router delle richieste (NEW\_ANALYSIS vs Q\_AND\_A)
- Coordinamento pipeline multi-agente
- Streaming status updates alla UI

#### Metodi Principali:

Metodo	Input	Output	Descrizione	
`determine_intent()`	query: str, has_context: bool	"NEW_ANALYSIS" \ "Q_AND_A"	Classifica intent via LLM o euristica	
`route_request()`	message, current_result, inputs	Generator[dict]	Entry point unificato	
`stream_workflow()`	request, inputs	Generator[dict]	Esegue pipeline 6-step	
`ask_question()`	query, context	Generator[dict]	Path Q&A	

#### Intent Detection Logic:

```
# LLM prompt
"Classify: NEW_ANALYSIS (new calculations) or Q_AND_A (question about existing)"

# Fallback heuristics
if "run" in query or "calculate" in query or "analyze" in query:
    return "NEW_ANALYSIS"
else:
    return "Q_AND_A"
```

#### Limiti:

- Intent detection dipende da LLM disponibile
- Nessun supporto per richieste ibride (analisi + domanda)

### 3.2 Methodology Agent

\*\*File:\*\* `src/agents/methodology.py`

\*\*Classe:\*\* `MethodologyAgent`

#### Responsabilità:

- Analizzare la richiesta utente
- Determinare il tipo di analisi (QUICK/STANDARD/FULL)
- Produrre configurazione per l'Execution Agent

**Input:** Stringa richiesta naturale

**\*\*Output:\*\*** `(ReservingConfigFile, AgentLog)`

### Logica di Mapping:

Keyword	Analysis Type	Methods Enabled
"quick", "fast"	QUICK	Chain Ladder only
"stress", "full", "complete"	FULL	All methods + stress testing
default	STANDARD	CL + Mack + Bootstrap

### Configurazione Risultante:

```
class ReservingConfigFile:  
    analysis_type: AnalysisType  
    run_model_selection: bool = True  
    run_bootstrap: bool = True  
    n_bootstrap_simulations: int = 1000  
    run_diagnostics: bool = True  
    run_stress_testing: bool = False # True solo per FULL
```

### Limiti:

- Mapping basato su keyword semplici
- Nessun parsing NLU avanzato
- Configurazione non parametrizzabile dall'utente

## 3.3 Selection Agent (LLM-Driven)

**\*\*File:\*\*** `src/agents/selection.py`

**\*\*Classe:\*\*** `SelectionAgent`

### Responsabilità:

- Calcolare TUTTI i 7 factor estimators
- Eseguire validazione completa (13 metriche)
- Eseguire 5 test diagnostici
- Analizzare volatilità e tail fitting
- Chiedere all'LLM di selezionare il metodo ottimale

### Metodi Principali:

Metodo	Descrizione
`analyze_and_select(triangle, premium)`	Entry point principale
`_calculate_all_estimators(triangle)`	Calcola 7 estimatori
`_run_full_validation(triangle)`	13 error metrics via ModelSelector
`_run_diagnostics(triangle)`	5 test Mack
`_analyze_volatility(triangle)`	CV, trend, structural breaks
`_fit_tail(triangle)`	7 curve fitting methods
`_calculate_maturity(triangle)`	% maturity per anno
`_calculate_bf_comparison(triangle, premium)`	BF vs CL
`_prepare_comprehensive_context(all_data)`	Formatta per LLM

Metodo	Descrizione
`_llm_select(context, all_data)`	Chiede all'LLM
`_fallback_selection(all_data)`	Regole se LLM non disponibile

## 7 Factor Estimators Calcolati:

1. Simple Average
2. Volume-Weighted
3. Geometric Average
4. Median
5. Leverage-Weighted
6. Exclude High/Low
7. Mack Adjusted

## LLM Selection Prompt:

Based on validation metrics, diagnostic tests, and maturity analysis, decide:

1. Which factor estimator to use as BASE method
2. Which years should use Bornhuetter-Ferguson (<70% maturity)
3. Whether to apply tail factor adjustment

RESPOND IN JSON:

```
{
  "selected_estimator": "...",
  "estimator_reason": "...",
  "bf_years": ["2022", "2023"],
  "bf_reason": "...",
  "summary": "..."
}
```

## Output:

```
{
  "selected_estimator": "Volume-Weighted",
  "estimator_reason": "Lowest MSE in out-of-sample validation",
  "bf_years": ["2022", "2023"],
  "bf_reason": "Years with <70% maturity",
  "all_estimators": {"Simple Average": 1234567, ...},
  "validation_metrics": {"Simple Average": {"MSE": 123, "MAE": 45}, ...},
  "maturity_by_year": {"2018": 95.2, "2019": 87.1, ...}
}
```

## Limiti:

- Dipende da LLM per selezione intelligente
- Fallback usa solo MSE (potrebbe non essere ottimale)
- Nessun supporto per preferenze utente

## 3.4 Execution Agent (Actuary)

\*\*File:\*\* `src/agents/reserving.py`

\*\*Classe:\*\* `ReservingExecutionAgent`

## Responsabilità:

- Caricare dati dal filesystem
- Eseguire calcoli attuariali via `EnhancedReservingWorkflow`

- Confezionare risultati in `ReservingOutput`

#### Metodi:

Metodo	Input	Output	
`execute(inputs, config)`	ReservingInput, ReservingConfigFile	ReservingOutput	
`_load_data(path, is_series)`	path: str	DataFrame \ Series	
`_package_results(workflow, config, triangle)`	...	ReservingOutput	

#### Pipeline di Esecuzione:

```

# 1. Load data
triangle = self._load_data(inputs.triangle_path)
premium = self._load_data(inputs.premium_path, is_series=True)

# 2. Create workflow
workflow = EnhancedReservingWorkflow(triangle, premium)

# 3. Run Chain Ladder (sempre)
workflow.run_chain_ladder()

# 4. Run Mack (se STANDARD o FULL)
if config.analysis_type != AnalysisType.QUICK:
    workflow.run_mack_model()

# 5. Run Bootstrap (se abilitato)
if config.run_bootstrap:
    workflow.run_bootstrap(n_simulations=config.n_bootstrap_simulations)

# 6. Run BF e Cape Cod (se FULL)
if config.analysis_type == AnalysisType.FULL:
    workflow.run_alternative_methods()

# 7. Run Diagnostics
if config.run_diagnostics:
    workflow.run_diagnostics()

# 8. Package results
return self._package_results(workflow, config, triangle)

```

#### Data Enrichment per LLM:

Il metodo `\_package\_results` popola:

- `triangle\_info`: metadata (n\_years, periods, range)
- `detailed\_data`: triangolo completo, fattori, riserve per anno
- Questi dati permettono all'LLM di rispondere a qualsiasi domanda

#### Limiti:

- Path hardcoded per alcuni file
- Nessun supporto per formati diversi da CSV
- Errori di caricamento non granulari

## 3.5 Validation Agent

\*\*File:\*\* `src/agents/validation.py`

\*\*Classe:\*\* `ValidationAgent`

**Responsabilità:**

- Verificare sanity dei risultati
- Confrontare metodi per coerenza
- Assegnare confidence score

**Controlli Eseguiti:**

Check	Condizione	Severity	Score Impact
Model Adequacy	score < 60	WARNING	-15
Mack CV	cv > 25%	WARNING	-10
CL vs Mack Divergence	diff > 10%	WARNING	-10
Bootstrap vs CL	diff > 15%	INFO	-5

**Calcolo Confidence Score:**

```

score = 100

if diagnostics.adequacy_score < 60:
    score -= 15
    issues.append(ValidationIssue(severity="WARNING", ...))

if mack.cv > 0.25:
    score -= 10
    issues.append(ValidationIssue(severity="WARNING", ...))

# etc.

return ValidationReport(
    status=ValidationStatus.WARNING if issues else ValidationStatus.PASSED,
    overall_confidence_score=max(0, score),
    issues=issues
)

```

**\*\*Output:\*\* `ValidationReport`**

**Limiti:**

- Soglie hardcoded
- Nessun check su BF o Cape Cod
- Manca confronto con benchmark industry

## 3.6 Reporting Agent

**\*\*File:\*\* `src/agents/reporting.py`**

**\*\*Classe:\*\* `ReportingAgent`**

**Responsabilità:**

- Generare report markdown completo
- Includere tutte le decisioni LLM
- Formattare tabelle e sezioni

**Sezioni Generate:**

- 1. Header:** Data, Status, Methodology
- 2. Data Overview:** Triangle metadata

### 3. ■ AI Method Selection:

- Selected estimator + reason
- BF years + reason
- All estimators comparison table
- Validation metrics table
- Maturity by year table
- AI Summary

### 4. Executive Summary: Reserve recommendation

### 5. Reserve Estimates by Method: Comparison table

### 6. Bootstrap Confidence Intervals: Percentiles

### 7. Development Factors: Factor table

### 8. Reserves by Accident Year: Year breakdown

### 9. Validation & Diagnostics: Findings

### 10. Footer: Disclaimer

**Output:** Markdown string (~3000-5000 chars)

### Limiti:

- Solo formato Markdown
- Nessun export PDF/Excel diretto
- Layout non customizzabile

## 3.7 QA Specialist Agent

**\*\*File:\*\* `src/agents/qa.py`**

**\*\*Classe:\*\* `QASpecialistAgent`**

### Responsabilità:

- Rispondere a domande sui risultati
- Usare LLM con contesto completo
- Fallback a pattern matching

### Strategia:

1. Se LLM disponibile:
  - Serializza ReservingOutput + ValidationReport in JSON
  - Passa come contesto all'LLM
  - LLM genera risposta naturale
2. Se LLM non disponibile:
  - Pattern matching su keyword
  - Estrae valori specifici dal contesto
  - Formatta risposta template

### Keyword Patterns (Fallback):

Pattern	Risposta
"validation", "issues"	Lista issues da ValidationReport
"ibnr", "reserve"	Total reserve da chain_ladder
"ultimate"	Ultimate loss
"mack", "error"	Standard error e CV
"bootstrap", "percentile"	Percentili distribuzione

Pattern	Risposta
"bornhuetter", "bf"	BF reserve
"cape cod"	Cape Cod reserve
"diagnostic"	Adequacy score e rating

### Limiti:

- Fallback molto limitato
- Nessun supporto per domande complesse
- Nessuna memoria conversazionale (oltre session state)

## 3.8 LLM Client

\*\*File:\*\* `src/agents/llm\_utils.py`

\*\*Classe:\*\* `LLMClient`

### Configurazione:

```
class LLMClient:
    def __init__(self):
        self.api_key = os.getenv("OPENAI_API_KEY")
        self.client = OpenAI(api_key=self.api_key) if self.api_key else None

    def is_available(self) -> bool:
        return self.client is not None

    def get_completion(self, system_prompt, user_prompt, model="gpt-4o-mini"):
        response = self.client.chat.completions.create(
            model=model,
            messages=[
                {"role": "system", "content": system_prompt},
                {"role": "user", "content": user_prompt}
            ],
            temperature=0.0 # Deterministic
        )
        return response.choices[0].message.content
```

\*\*Modello Default:\*\* `gpt-4o-mini` (veloce, economico)

**Graceful Degradation:** Tutti i chiamanti hanno logica fallback se LLM non disponibile.

## 4. Reserving Package Integration

### 4.1 Chiamate al Package

Il sistema usa internamente queste componenti:

Component	File	Chiamato da
`ChainLadder`	`chain_ladder.py`	ExecutionAgent, SelectionAgent
`MackChainLadder`	`stochastic_reserving/mack_model.py`	ExecutionAgent

Component	File	Chiamato da
`ODPBootstrap`	`stochastic_reserving/bootstrap.py`	ExecutionAgent
`BornhuetterFerguson`	`alternative_methods/bornhuetter_ferguson.py`	ExecutionAgent, SelectionAgent
`CapeCod`	`alternative_methods/cape_cod.py`	ExecutionAgent, SelectionAgent
`get_all_estimators()`	`model_selection/factor_estimators.py`	SelectionAgent
`ModelSelector`	`model_selection/model_selector.py`	SelectionAgent
`DiagnosticTests`	`diagnostics/diagnostic_tests.py`	SelectionAgent
`VolatilityAnalyzer`	`diagnostics/volatility_analysis.py`	SelectionAgent
`TailEstimator`	`tail_fitting/tail_estimator.py`	SelectionAgent

## 4.2 Flow dei Parametri

```
Triangle (DataFrame)
    [
        ChainLadder(triangle)
            .calculate_age_to_age_factors()
            .select_development_factors()
            .calculate_cumulative_factors()
            .project_ultimate_losses()

        MackChainLadder(triangle)
            .fit()
            .get_confidence_intervals()

        ODPBootstrap(triangle, n_simulations=1000)
            .fit()
            .reserve_distribution

        BornhuetterFerguson(triangle, premium)
            .fit()
            .get_comparison()

        get_all_estimators()
            [est.estimate(triangle) for est in estimators]
```

## 4.3 Gestione Risultati

Ogni metodo produce output specifico che viene mappato in Pydantic models:

```
# Chain Ladder → MethodResult
MethodResult(
    method_name="Chain Ladder",
    total_reserve=cl.ultimate_losses['Reserve'].sum(),
    ultimate_loss=cl.ultimate_losses['Ultimate'].sum()
)

# Mack → StochasticResult
StochasticResult(
    method_name="Mack",
    total_reserve=...,
    standard_error=mack.mse_reserves['SE'].sum(),
    cv=mack.mse_reserves['CV'].mean(),
    percentiles={"75%": ..., "95%": ...}
)
```

## 5. Data & Artifacts

### 5.1 Pydantic Models (*schemas.py*)

#### Enums:

Enum	Values	Usage
`AgentRole`	ORCHESTRATOR, METHODOLOGY, EXECUTION, VALIDATION, REPORTING	Audit trail
`AnalysisType`	QUICK, STANDARD, FULL	Config
`ValidationStatus`	PASSED, WARNING, REJECTED	Validation

#### Core Models:

```
ReservingInput
    triangle_path: str
    premium_path: Optional[str]
    loss_ratios_path: Optional[str]

ReservingConfigFile
    analysis_type: AnalysisType
    run_model_selection: bool
    run_bootstrap: bool
    n_bootstrap_simulations: int
    run_diagnostics: bool
    run_stress_testing: bool

ReservingOutput
    timestamp: datetime
    config_used: ReservingConfigFile
    triangle_info: TriangleMetadata
    detailed_data: DetailedData
    method_selection: MethodSelection
    chain_ladder: MethodResult
    mack: StochasticResult
    bootstrap: StochasticResult
    bornhuetter_ferguson: MethodResult
    cape_cod: MethodResult
    diagnostics: DiagnosticsResult

ValidationReport
    status: ValidationStatus
    overall_confidence_score: int (0-100)
    issues: List[ValidationIssue]
    comparison_summary: str

AgentLog
    timestamp: datetime
    agent: AgentRole
    action: str
    details: str
```

### 5.2 Storage Locations

Artifact	Location	Format
Input Triangle	`data/processed/reported_absolute_losses.csv`	CSV
Premium	`data/processed/earned_premium.csv`	CSV
Final Report	`outputs/agent_runs/final_report.md`	Markdown

Artifact	Location	Format
Session State	In-memory (Streamlit)	Python dict

## 5.3 Audit Trail

Ogni agente produce un `AgentLog`:

```
AgentLog(
    timestamp=datetime.now(),
    agent=AgentRole.EXECUTION,
    action="Run Chain Ladder",
    details="Completed with reserve = $1,234,567"
)
```

I log vengono aggregati in `audit\_trail: List[AgentLog]` e mostrati nella UI.

# 6. Question & Re-Analysis Loop

## 6.1 Intent Detection

```
def determine_intent(self, query: str, has_context: bool) -> str:
    # 1. Try LLM
    if self.llm.is_available():
        prompt = f"Classify: NEW_ANALYSIS or Q_AND_A\nQuery: {query}"
        response = self.llm.get_completion(system_prompt, prompt)
        if "NEW_ANALYSIS" in response:
            return "NEW_ANALYSIS"
        return "Q_AND_A"

    # 2. Fallback heuristics
    analysis_keywords = ["run", "calculate", "analyze", "execute", "perform"]
    if any(kw in query.lower() for kw in analysis_keywords):
        return "NEW_ANALYSIS"
    return "Q_AND_A"
```

## 6.2 Q&A; Flow (Nessun Nuovo Calcolo)

```
User: "What is the CV for the Mack model?"
■
▼
Orchestrator.determine_intent() → "Q_AND_A"
■
▼
QASpecialistAgent.answer_query(query, context)
■
■■■ context = st.session_state.final_result
■
■■■ LLM: "Based on the Mack model results, the CV is 18.5%..."
■ OR
■■■ Fallback: "The Mack model shows CV: 18.5%"
■
▼
Response displayed in chat
```

## 6.3 Re-Analysis Flow (Nuovi Calcoli)

```
User: "Run a full analysis with stress testing"
└─
    └─
        Orchestrator.determine_intent() → "NEW_ANALYSIS"
        └─
            └─
                stream_workflow()
                └─
                    ┌─ Methodology: analysis_type=FULL, run_stress_testing=True
                    ┌─ Selection: Calcola tutti estimatori
                    ┌─ Execution: Run tutti i metodi
                    ┌─ Validation: Check risultati
                    ┌─ Reporting: Genera report
                    ┌─ Synthesis: Risposta finale
                    └─
                st.session_state.final_result = new_results
```

## 6.4 Limitazioni

- **Nessuna Memoria Conversazionale:** Ogni domanda è indipendente (solo session state)
- **Nessun Calcolo On-Demand:** Q&A non può triggerare calcoli specifici
- **Intent Binario:** Solo NEW\_ANALYSIS o Q\_AND\_A, niente di intermedio

# 7. Example Execution Paths

## 7.1 Path 1: Analisi Standard

**User Input:** "Analyze the triangle"

**Files Coinvolti:**

```
pages/3_■_AI_Team.py
    ┌─ Orchestrator.route_request()
        ┌─ determine_intent() → "NEW_ANALYSIS"
        ┌─ stream_workflow()
            ┌─ methodology.py → MethodologyAgent.plan_analysis()
            ┌─     ┌─ Returns: AnalysisType.STANDARD
            ┌─
            ┌─ selection.py → SelectionAgent.analyze_and_select()
            ┌─     ┌─ model_selection/factor_estimators.py
            ┌─     ┌─ model_selection/model_selector.py
            ┌─     ┌─ diagnostics/diagnostic_tests.py
            ┌─     ┌─ Returns: {"selected_estimator": "Volume-Weighted", ...}
            ┌─
            ┌─ reserving.py → ReservingExecutionAgent.execute()
            ┌─     ┌─ data_loader.py
            ┌─     ┌─ chain_ladder.py
            ┌─     ┌─ stochastic_reserving/mack_model.py
            ┌─     ┌─ stochastic_reserving/bootstrap.py
            ┌─     ┌─ Returns: ReservingOutput
            ┌─
            ┌─ validation.py → ValidationAgent.validate()
            ┌─     ┌─ Returns: ValidationReport
            ┌─
            ┌─ reporting.py → ReportingAgent.generate_report()
```

■■■ Returns: Markdown string

**Output:** Report completo con CL, Mack, Bootstrap, 7 estimatori confrontati.

## 7.2 Path 2: Analisi Full (con Stress Testing)

**User Input:** "Run full analysis with stress testing"

**Differenze da Standard:**

```
methodology.py
    ■■■ Keywords "full", "stress" detected
    ■■■ Returns: AnalysisType.FULL, run_stress_testing=True

reserving.py
    ■■■ Additional calls:
        ■■■ alternative_methods/bornhuetter_ferguson.py
        ■■■ alternative_methods/cape_cod.py
        ■■■ scenario_analysis/stress_testing.py
```

**Output Aggiuntivo:** BF reserve, Cape Cod reserve, stress scenarios.

## 7.3 Path 3: Domanda Post-Analisi

**User Input:** "How many accident years are in the triangle?"

\*\*Pre-condizione:\*\* `st.session\_state.final\_result` esiste

**Files Coinvolti:**

```
pages/3_■_AI_Team.py
    ■■■ Orchestrator.route_request()
        ■■■■ determine_intent() → "Q_AND_A"
        ■■■■ ask_question()
            ■■■■ qa.py → QASpecialistAgent.answer_query()
                ■■■■■ context = final_result["structured_results"]
                ■■■■■ LLM available:
                    ■■■■■■■ "The triangle contains 10 accident years (2014-2023)"
                ■■■■■ LLM not available:
                    ■■■■■■■ Keyword match on "years" → extract from triangle_info
```

**Output:** Risposta testuale senza nuovi calcoli.

# 8. Gaps & Improvement Opportunities

## 8.1 Architetturali

Gap	Impatto	Soluzione Proposta
**Nessun Code Interpreter**	Q&A non può fare calcoli custom	Aggiungere CodeAgent con sandbox Python
**Intent Detection Binario**	Richieste ibride mal gestite	Classificazione multi-label o chain-of-thought
**Nessuna Memoria Conversazionale**	Ogni domanda isolata	Aggiungere conversation buffer o RAG
**Soglie Hardcoded**	Validazione non configurabile	Esternalizzare in config YAML

## 8.2 Calcoli

Gap	Impatto	Soluzione Proposta
**Selection Agent Overkill**	Calcola TUTTO ogni volta	Lazy evaluation o caching per estimatori
**Bootstrap Lento**	1000 simulazioni bloccanti	Async/background job o progressive loading
**Nessun Caching Risultati**	Re-calcio su ogni run	Redis/pickle per risultati intermedi

## 8.3 UI/UX

Gap	Impatto	Soluzione Proposta
**Solo Markdown Report**	Nessun export professionale	Aggiungere PDF via ReportLab o FPDF
**Grafici Non nel Report**	Separazione report/charts	Integrare Plotly export in report
**Nessun Upload Custom**	Solo triangoli predefiniti	Aggiungere file uploader con validazione

## 8.4 Robustezza

Gap	Impatto	Soluzione Proposta
**LLM Single Point of Failure**	Fallback limitati	Migliorare fallback con regole più ricche
**Nessun Rate Limiting**	OpenAI 429 errors	Implementare retry con exponential backoff
**Errori Non Granulari**	Debug difficile	Structured logging con livelli

## 8.5 Testing

Gap	Impatto	Soluzione Proposta
**Cartella tests/ Vuota**	Nessun test automatizzato	Unit test per ogni agent
**Nessun Integration Test**	Regressioni non catturate	End-to-end test con fixtures

## Appendix: Quick Reference

## Agent Responsibilities

ORCHESTRATOR	███████	Router, Coordinator
METHODOLOGY	███	Config Planning (QUICK/STANDARD/FULL)
SELECTION	███	LLM Method Selection (7 estimators)
EXECUTION	███	Actuarial Calculations
VALIDATION	███	Sanity Checks, Scoring
REPORTING	███	Markdown Generation
QA SPECIALIST	███	Question Answering

## Data Flow

CSV Files → DataFrame → Calculations → Pydantic Models → JSON/Markdown → UI

## Key Files to Modify

Per...	Modifica...
Aggiungere metodo	`reserving.py`, `schemas.py`
Nuova validazione	`validation.py`
Cambiare report	`reporting.py`
Nuovi chart	`pages/3_■_AI_Team.py`
Nuovi test diagnostici	`selection.py`
Cambiare LLM	`llm_utils.py`

\*Documento generato automaticamente dall'analisi del codice sorgente.\*